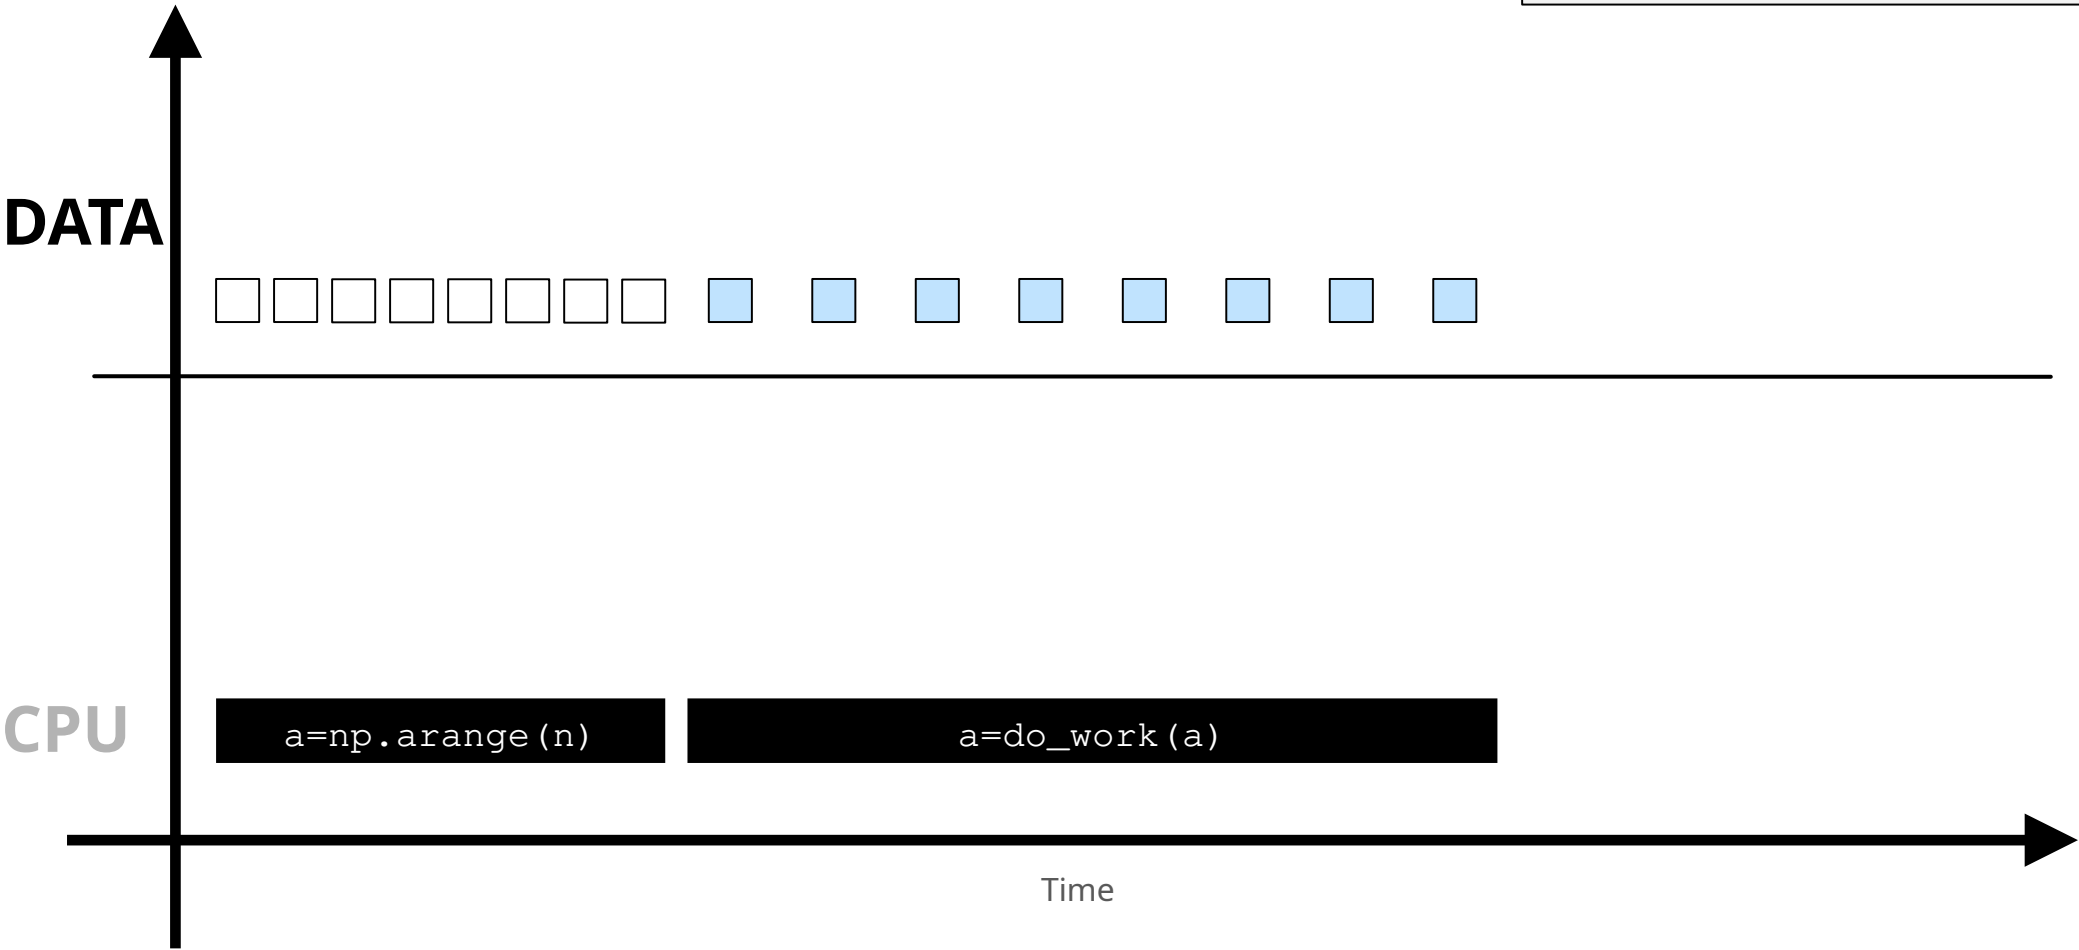


GPU-accelerated vs. CPU-only Applications

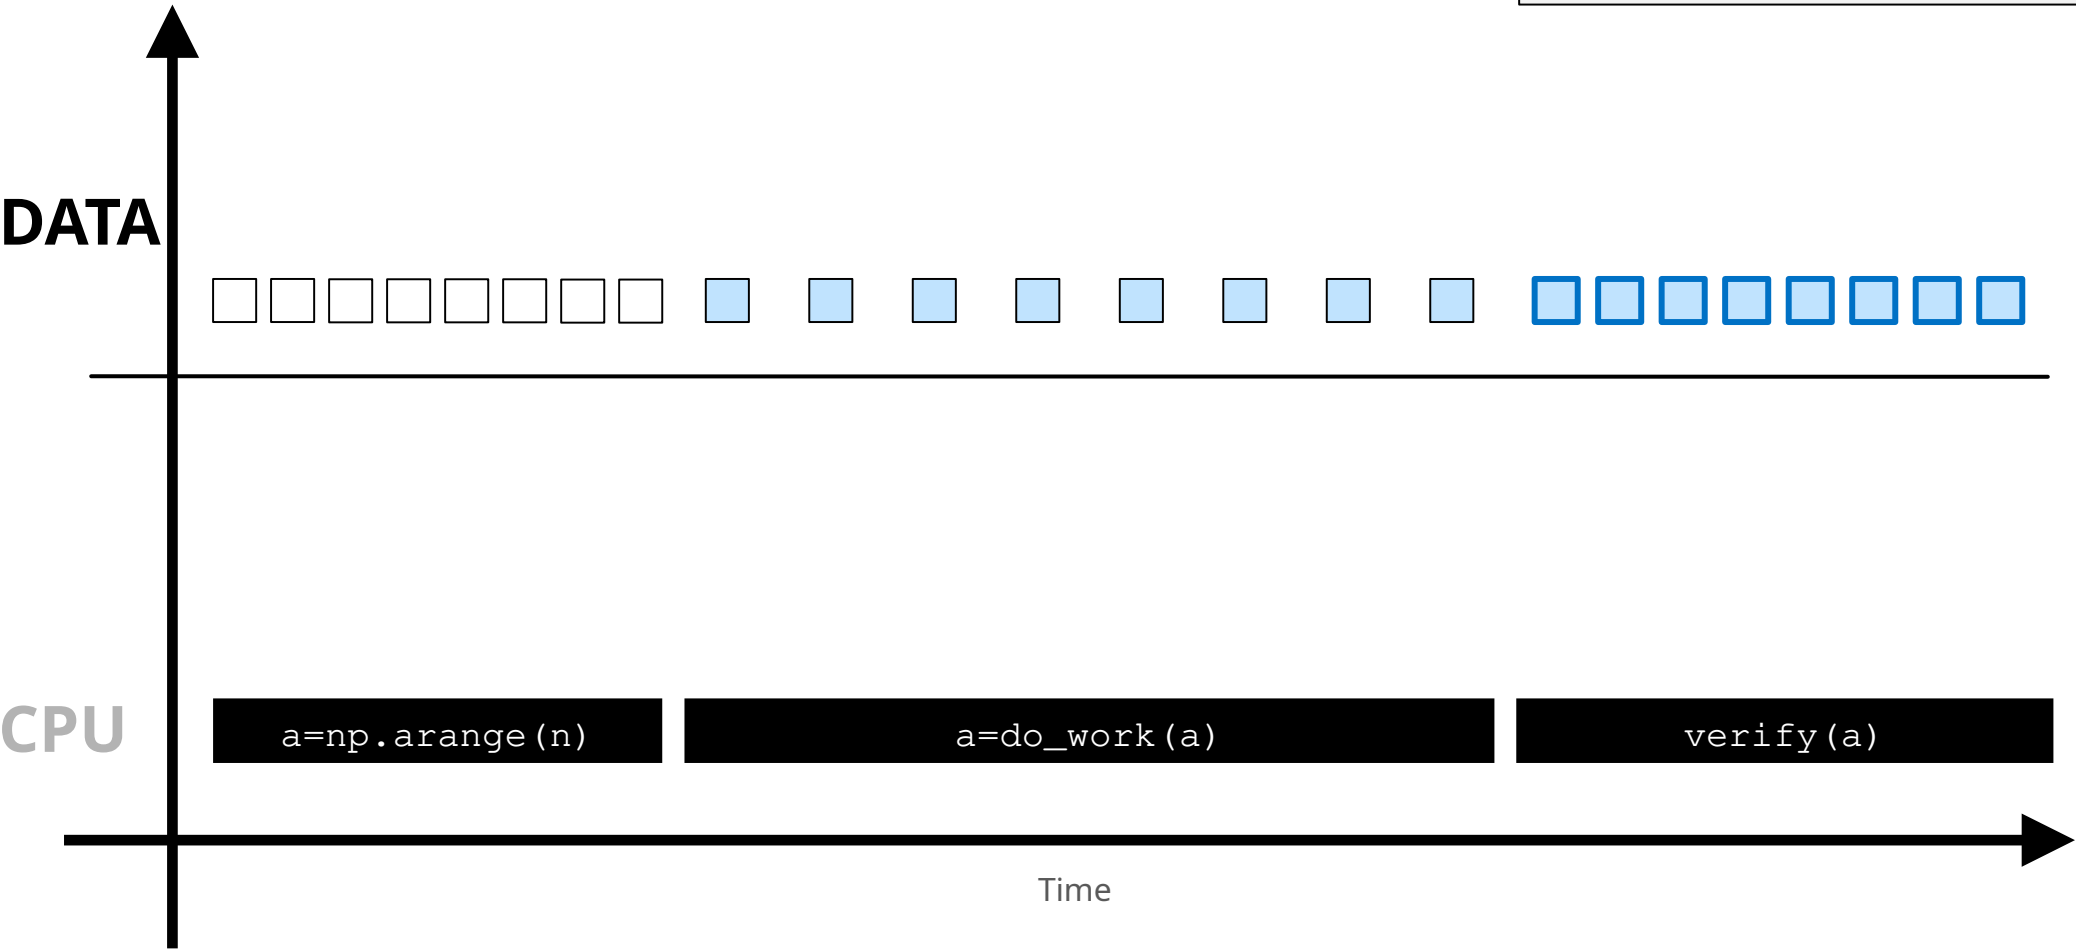
In **CPU-only applications** data is allocated on the CPU



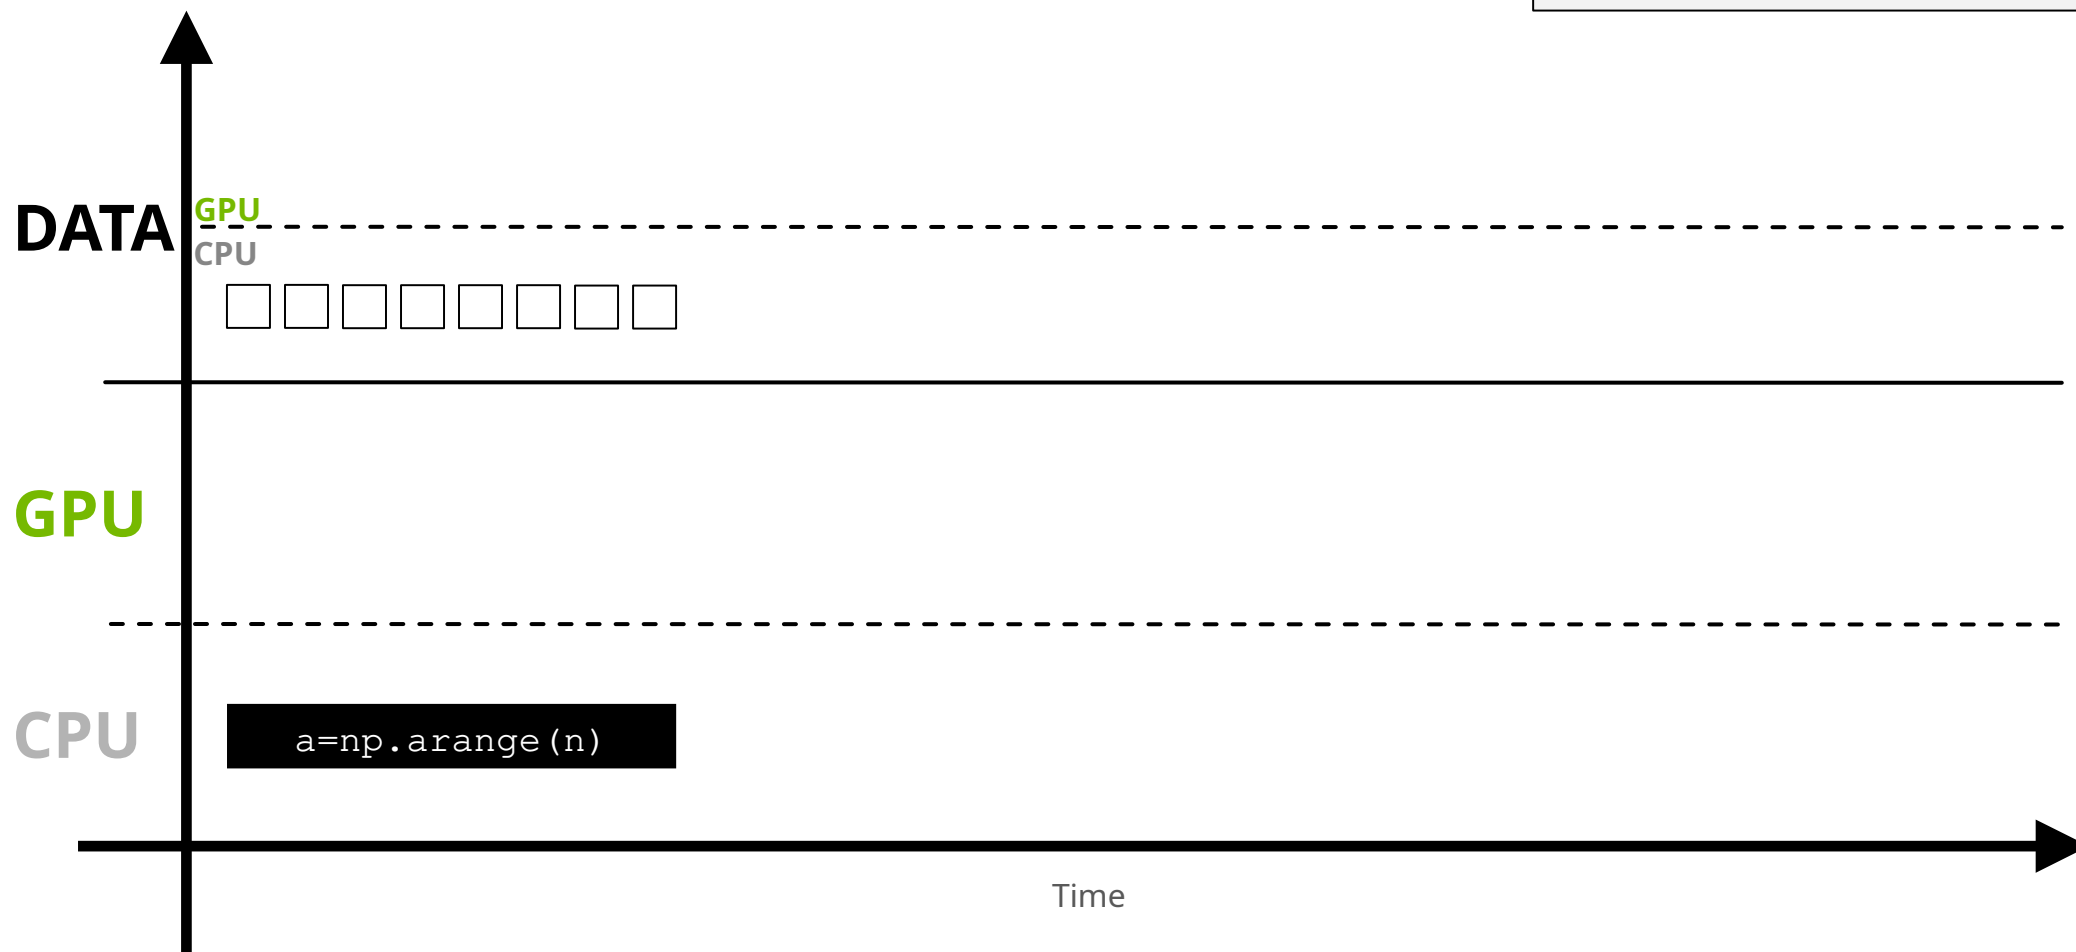
...and all work is performed serially on the CPU



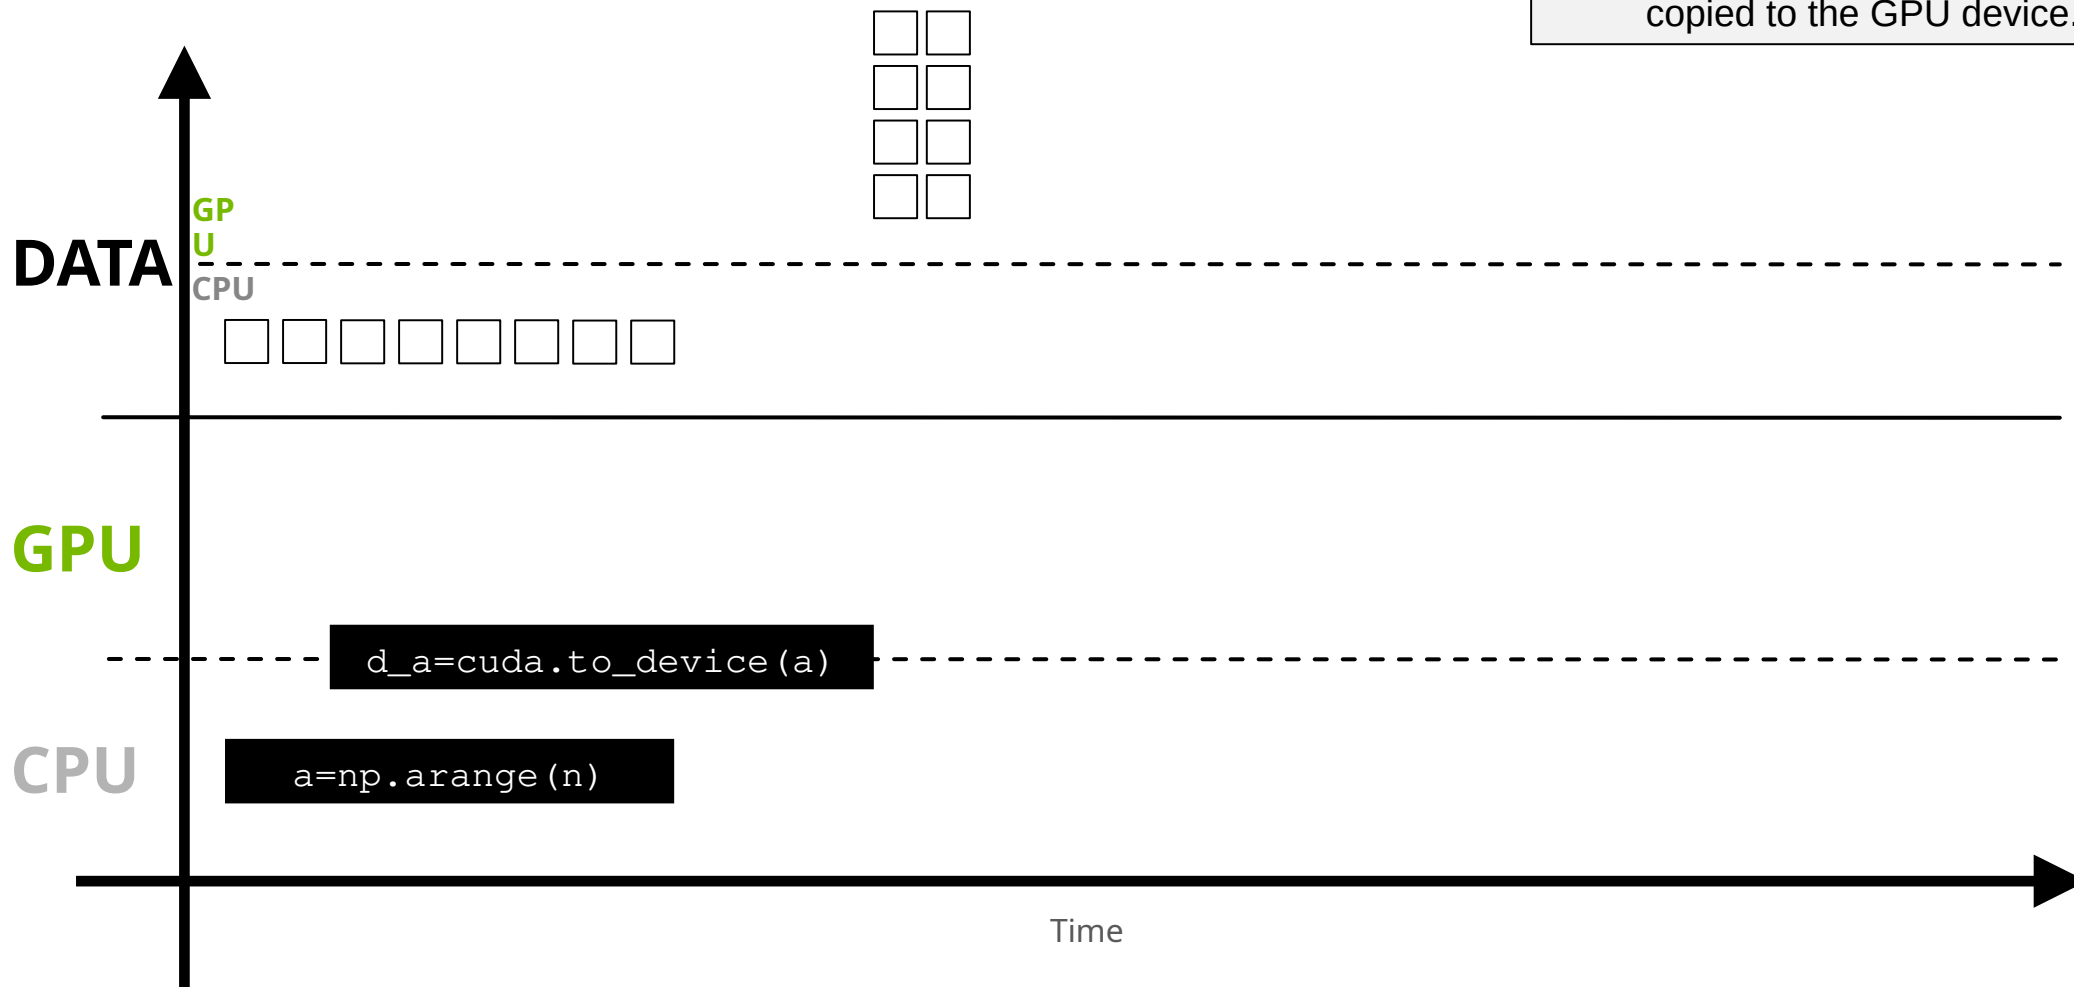
...and all work is performed serially on the CPU



In **accelerated applications** there is both host and device memory.




Data initialized on the CPU can be copied to the GPU device...



DATA

GPU
CPU



... where it can be worked on in parallel

GPU

```
d_a=do_work(d_a)
```

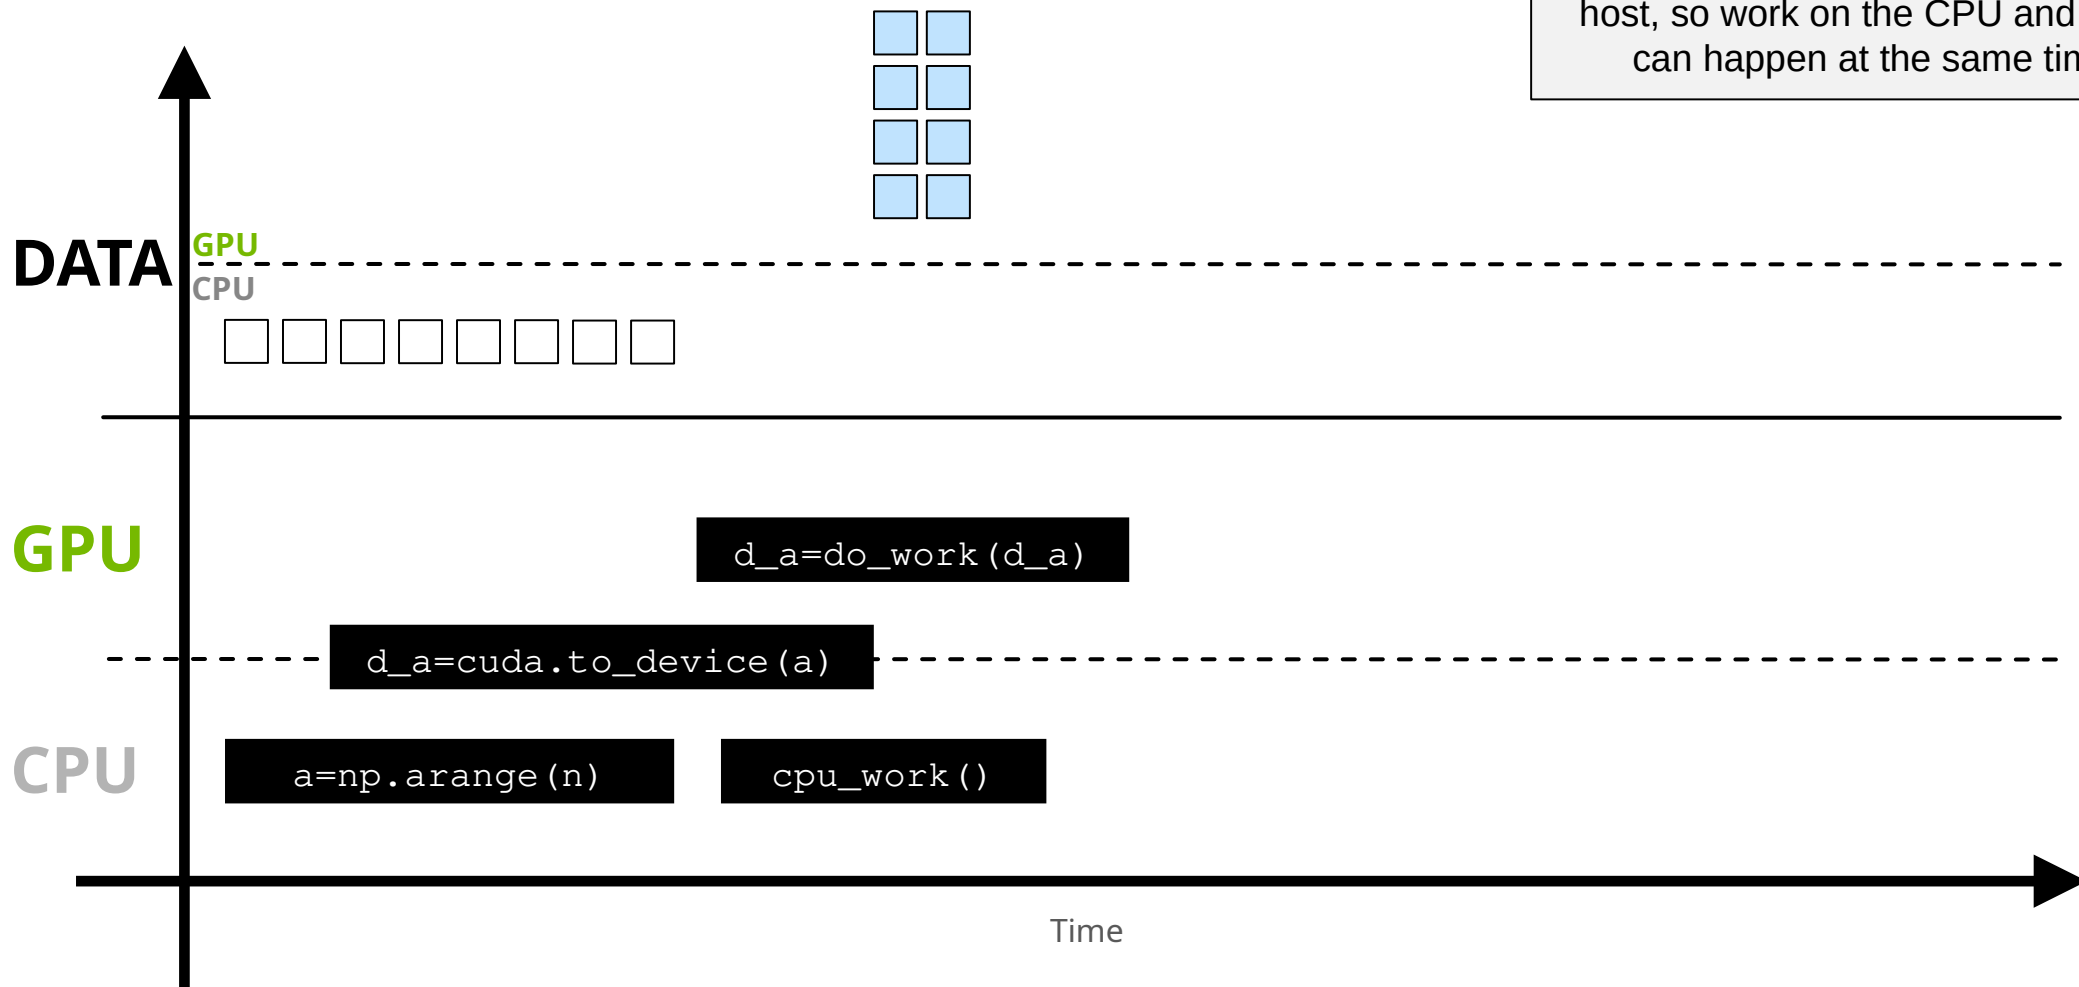
```
d_a=cuda.to_device(a)
```

CPU

```
a=np.arange(n)
```


Time

GPU work is asynchronous to the host, so work on the CPU and GPU can happen at the same time



DATA

GPU
CPU



GPU

`d_a=cuda.to_device(a)`

`d_a=do_work(d_a)`

CPU

`a=np.arange(n)`

`cpu_work()`


synchronize

Time

Programmers can indicate
synchronization points between the
CPU and GPU with
`cuda.synchronize()`

DATA

GPU
CPU



And data can be copied back to the CPU...

GPU

`d_a=do_work(d_a)`

`d_a=cuda.to_device(a)`

`a=d_a.copy_to_host()`

CPU

`a=np.arange(n)`


`cpu_work()`

synchronize

Time

DATA

GPU
CPU



... for verification etc.

GPU

```
d_a=cuda.to_device(a)
```

```
d_a=do_work(d_a)
```

```
a=d_a.copy_to_host()
```

CPU

```
a=np.arange(n)
```

```
cpu_work()
```

```
verify(a)
```

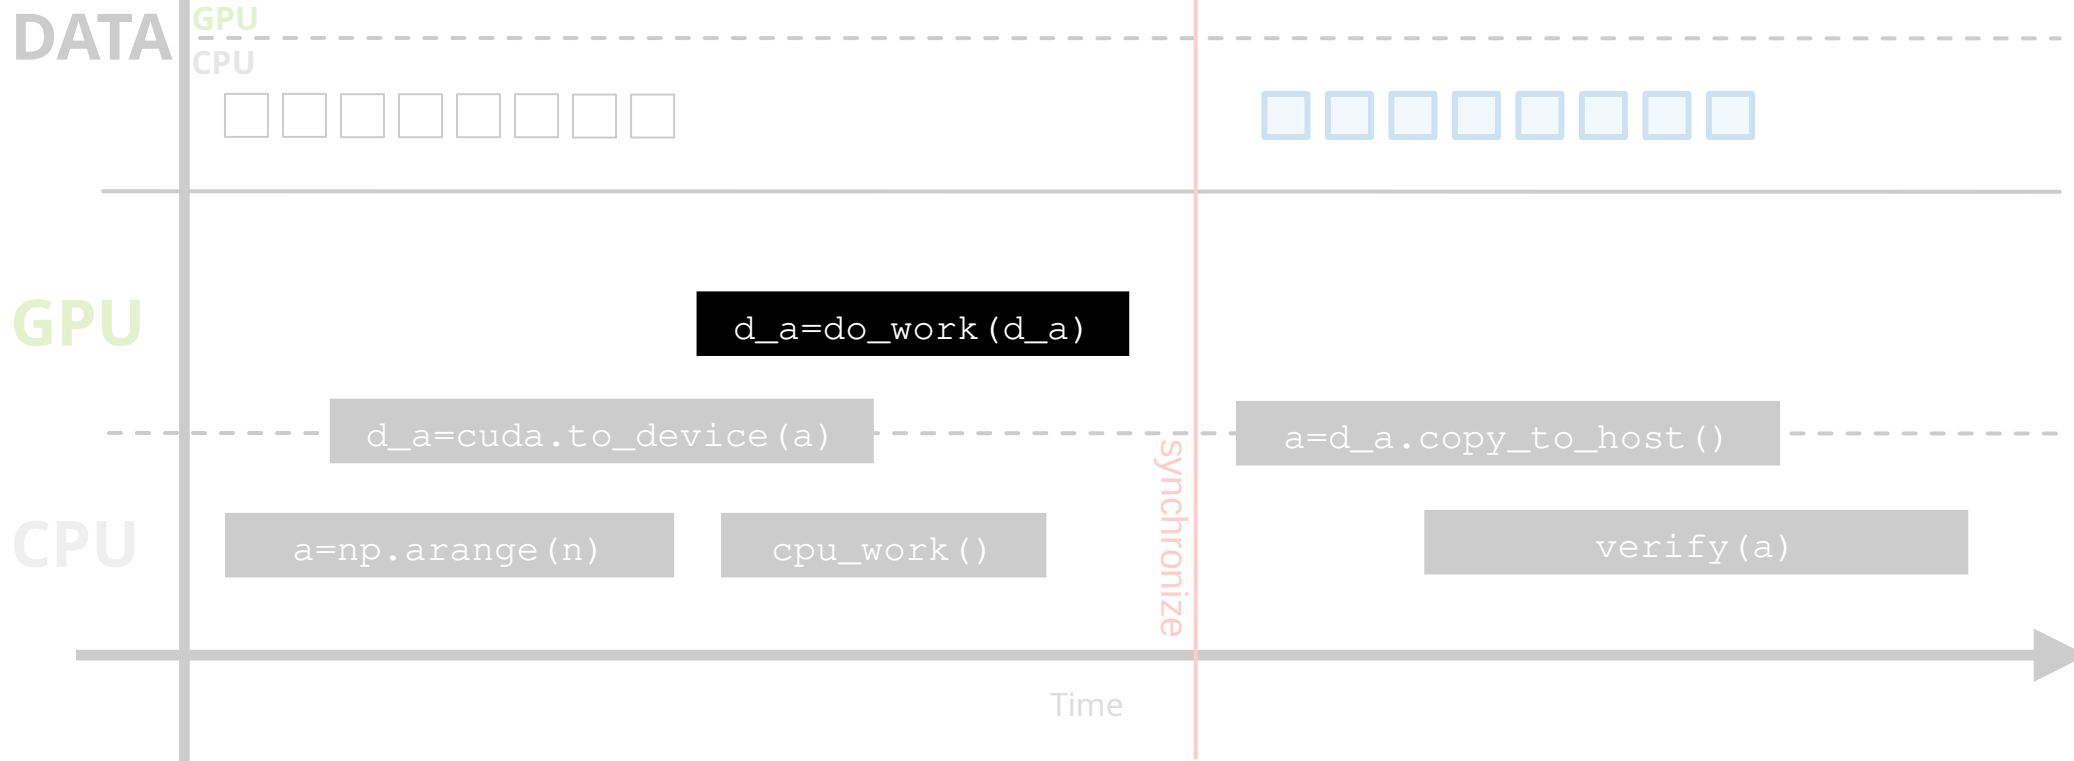
synchronize

Time


CUDA Thread Hierarchy

The background of the slide features a smooth gradient from a deep blue at the top to a vibrant green at the bottom. Overlaid on this gradient is a complex, abstract network of thin white lines and small white dots, resembling a neural network or a data visualization of interconnected nodes.


Let's dig into what happens when we launch a function on the GPU



GPU




GPU




Many threads run in parallel

GPU




CUDA can process thousands of threads in parallel. The sizes are greatly reduced in these images for simplicity.

GPU




A collection of threads is a **block**

GPU




There can be many blocks

GPU




A collection of blocks associated with a given kernel launch is a **grid**

GPU




GPU




Kernels are **launched** with an
execution configuration

GPU




The execution configuration defines
the number of blocks in the grid

GPU




... as well as the number of threads in each block

GPU



Every block in the grid contains the same number of threads


GPU



CUDA-Provided Thread Hierarchy Variables


Inside kernel definitions, CUDA-provided variables describe its executing thread, block, and grid

GPU



`gridDim.x` is the number of blocks in the grid, in this case 2


GPU



2


blockIdx.x is the index of the current block within the grid, in this case 0

GPU




blockIdx.x is the index of the current block within the grid, in this case **1**

GPU



Inside a kernel `blockDim.x`
describes the number of threads in a
block. In this case **4**


GPU



4


All blocks in a grid contain the same number of threads

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 0

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 1

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 2

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 3

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 0

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 1

GPU




Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 2

GPU



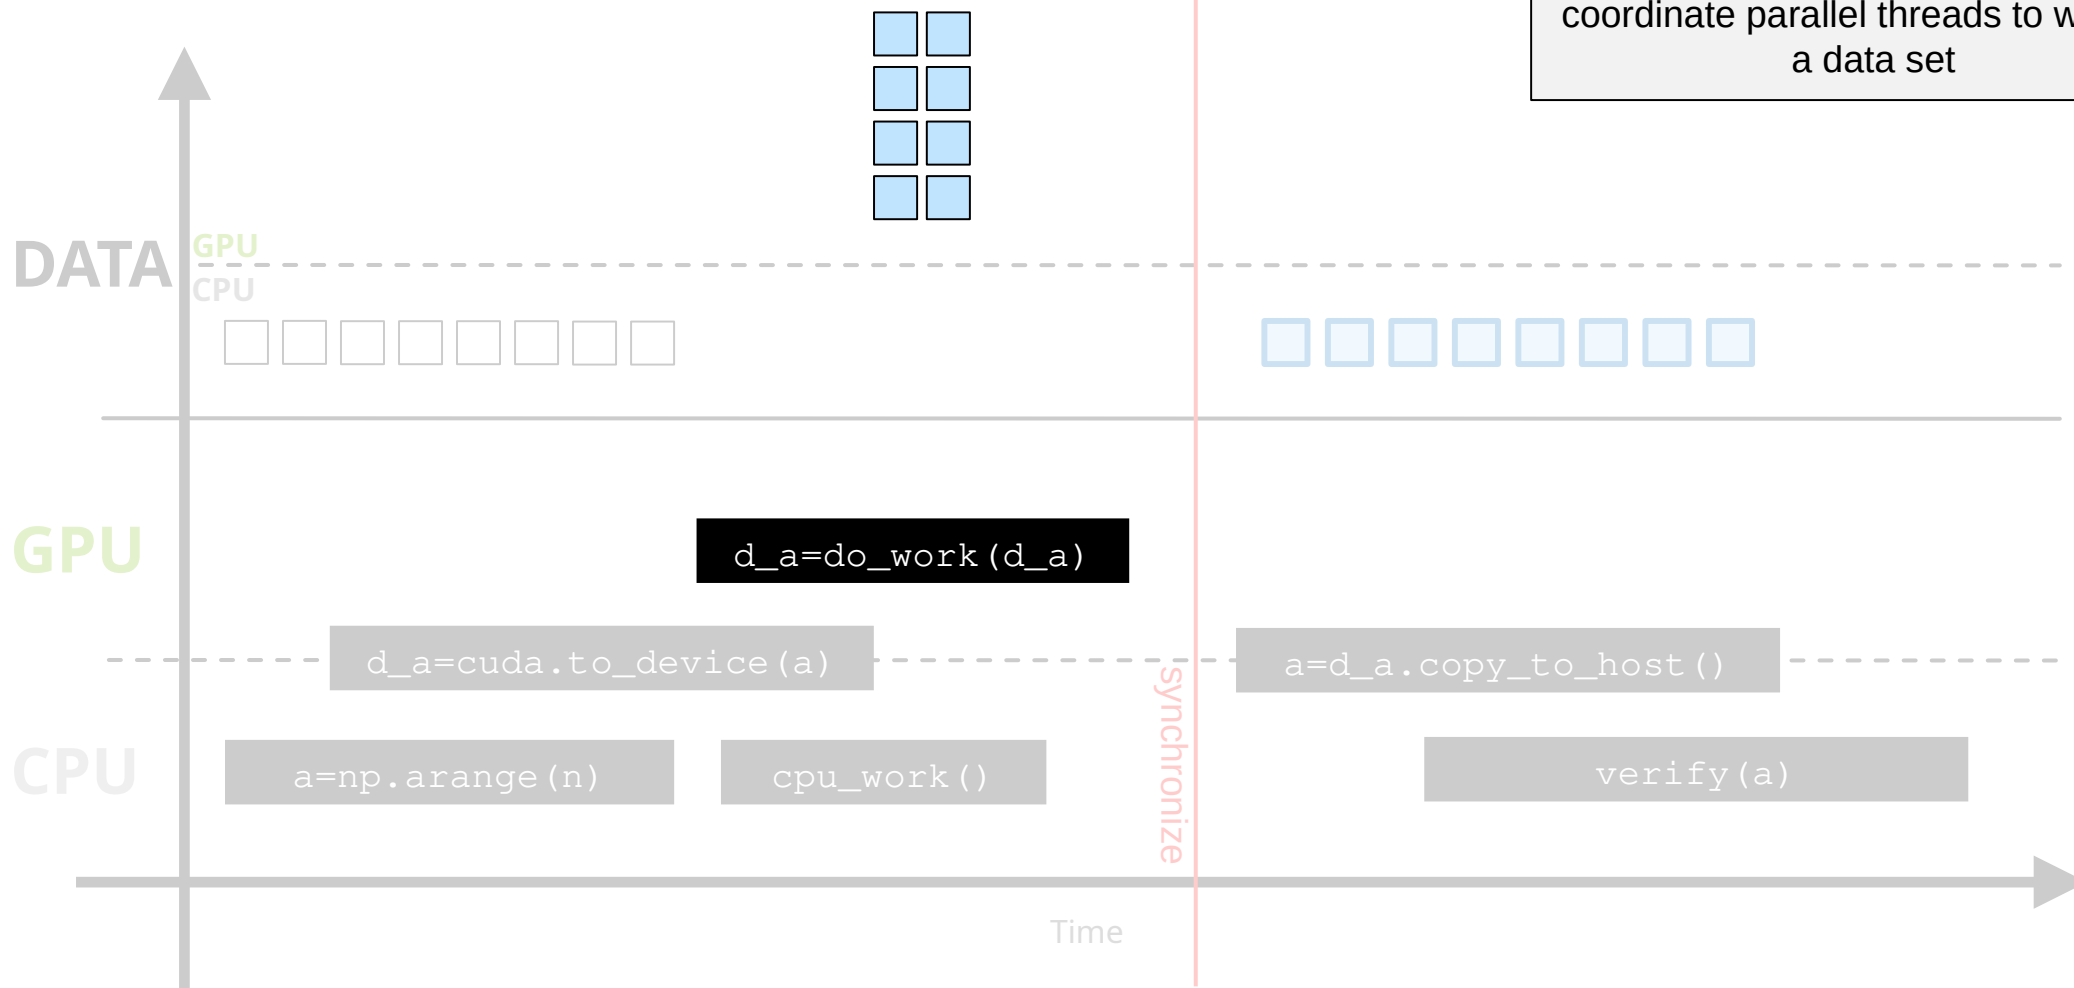
Inside a kernel `threadIdx.x`
describes the index of the thread
within a block. In this case 3

GPU




Coordinating Parallel Threads

Let's look at some basic ways to coordinate parallel threads to work on a data set




GPU
DATA



Assume data is in a 0 indexed vector

GPU




GPU DATA

0	4
1	5
2	6
3	7

Assume data is in a 0 indexed vector

GPU

```
do_work[2, 4] (d_a)
```




GPU
DATA


0	4
1	5
2	6
3	7

Somehow, each thread must be mapped to work on elements in the data

GPU




GPU
DATA




... if we can calculate a thread's index
within the entire grid, then we could
map that index to an index in the data

GPU




GPU
DATA



... unfortunately CUDA does not provide a single variable to capture this, only thread indices *within the block*

GPU




GPU DATA

0	4
1	5
2	6
3	7

There is an idiomatic way to calculate this value, however. Recall that each thread has access to the size of its block via `blockDim.x`

GPU




GPU DATA

0	4
1	5
2	6
3	7

...and the index of its block within the grid via `blockIdx.x`

GPU




GPU DATA

0	4
1	5
2	6
3	7

...and its own index within its block via
`threadIdx.x`

GPU




GPU DATA

0	4
1	5
2	6
3	7

Using these variables, the formula **`threadIdx.x + blockIdx.x * blockDim.x`** will return the thread's unique index in the whole grid, which we can then map to data elements.

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
0		0		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
0		0		4
data_index				
0				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
1		0		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
1		0		4
data_index				
1				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
2		0		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
2		0		4
data_index				
2				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
3		0		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
3		0		4
data_index				
3				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
0		1		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
0		1		4
data_index				
4				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
1		1		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
1		1		4
data_index				
5				

GPU




GPU DATA

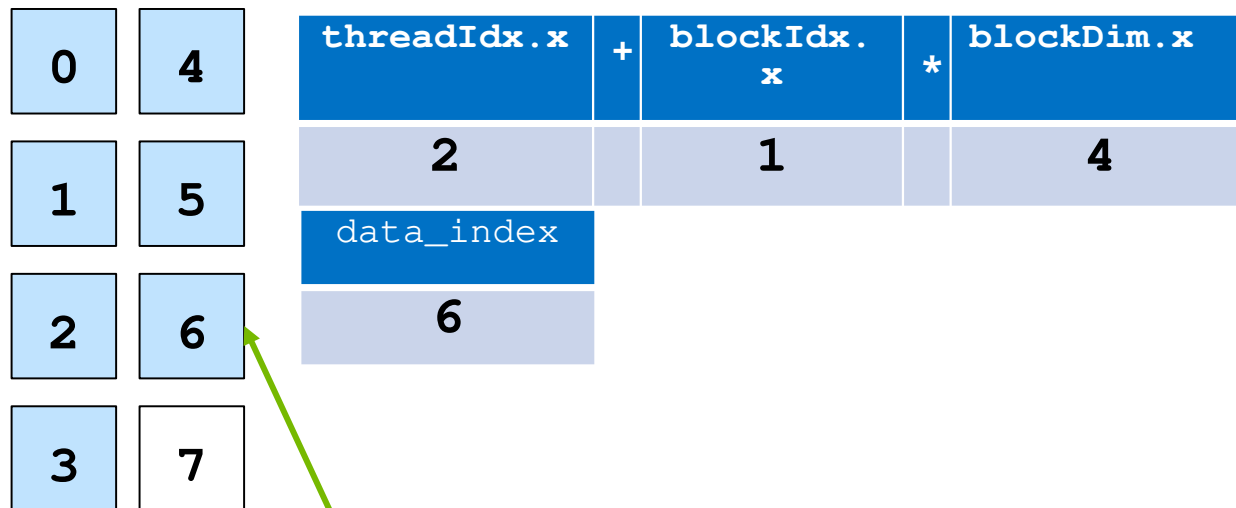
0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
2		1		4
data_index				
?				


GPU



GPU DATA



GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
3		1		4
data_index				
?				

GPU




GPU DATA

0	4
1	5
2	6
3	7

threadIdx.x	+	blockIdx.x	*	blockDim.x
3		1		4
data_index				
7				

GPU




GPU DATA

0	4	<code>threadIdx.x</code>	+	<code>blockIdx.x</code>	*	<code>blockDim.x</code>
1	5	3		1		4
2	6	<code>data_index</code>				
3	7	7				
		<code>grid(1)</code>				
		7				

As a convenience, Numba provides the `cuda.grid()` function, which will return a thread's unique index in the grid.

GPU





DEEP
LEARNING
INSTITUTE

www.nvidia.com/dli